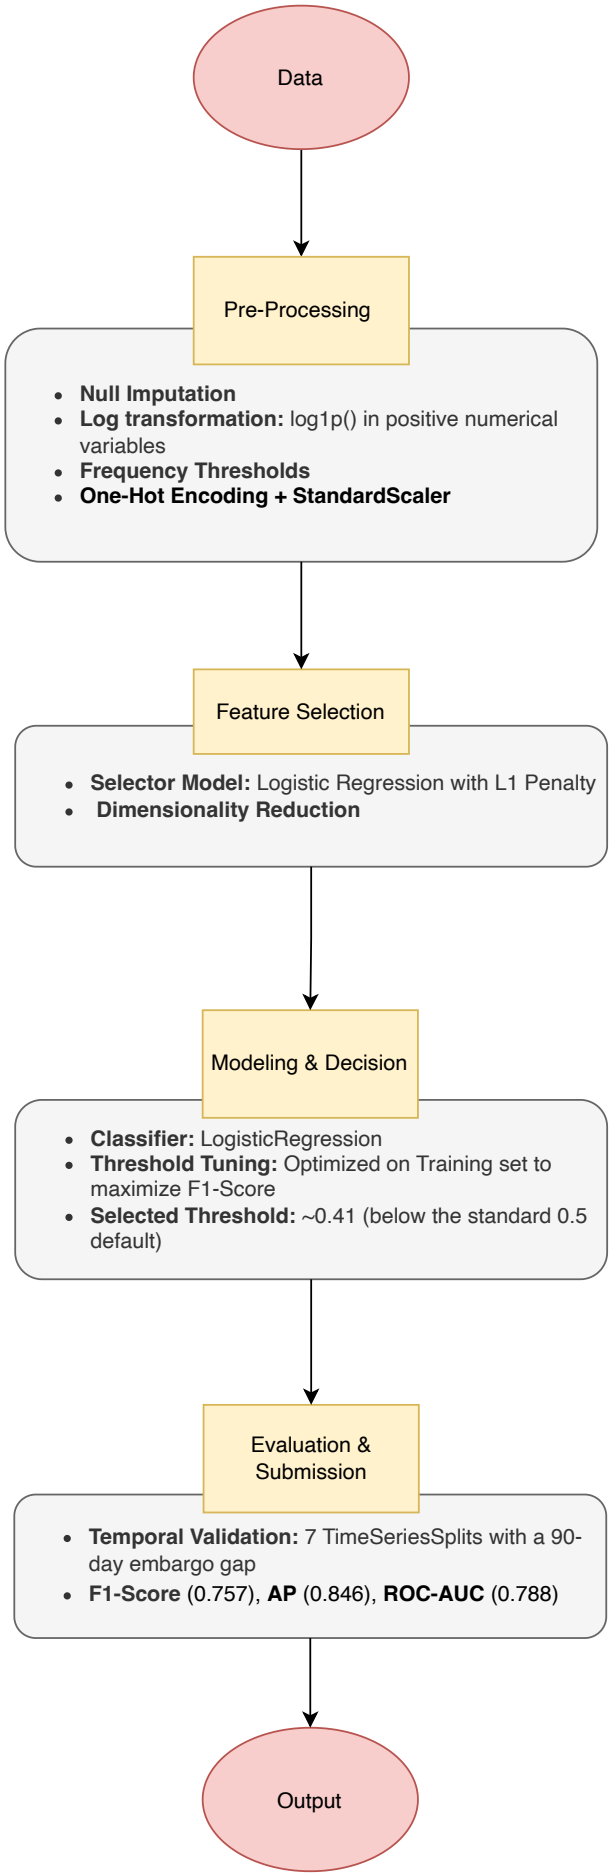
Data
Pre-Processing
Null Imputation
Log transformation: log1p() in positive numerical variables
Frequency Thresholds
One-Hot Encoding + StandardScaler
Feature Selection
Selector Model: Logistic Regression with L1 Penalty
Dimensionality Reduction
Modeling & Decision
Classifier: LogisticRegression
Threshold Tuning: Optimized on Training set to maximize F1-Score
Selected Threshold: ~0.41 (below the standard 0.5 default)
Evaluation & Submission
Temporal Validation: 7 TimeSeriesSplits with a 90-day embargo gap
F1-Score (0.757), AP (0.846), ROC-AUC (0.788)
Output